

```
import numpy as np
import pandas as pd
import nltk
import matplotlib.pyplot as plt
import seaborn as sbn
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
import re
from sklearn.feature_extraction.text import TfidfVectorizer
nltk.download('stopwords')
%matplotlib inline
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
## load dataset
df = pd.read_csv("bbc-text.csv")
df.head(10)
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

	category	text
0	tech	tv future in the hands of viewers with home th...
1	business	worldcom boss left books alone former worldc...
2	sport	tigers wary of farrell gamble leicester say ...
3	sport	yeading face newcastle in fa cup premiership s...
4	entertainment	ocean s twelve raids box office ocean s twelve...
5	politics	howard hits back at mongrel jibe michael howar...
6	politics	blair prepares to name poll date tony blair is...
7	sport	henman hopes ended in dubai third seed tim hen...
8	sport	wilkinson fit to face edinburgh england captai...
9	entertainment	last star wars not for children the sixth an...

```
df.shape
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
(2225, 2)
```

```
df['text'][0]
```

```
↳ 'tv future in the hands of viewers with home theatre systems plasma high-definition tvs and digital video recorders moving into th
```

```
df['category'].unique()
```


```
↳ array(['tech', 'business', 'sport', 'entertainment', 'politics'],  
      dtype=object)
```

```
df['category'].value_counts()
```

```
↳ sport      511  
   business   510  
   politics   417  
   tech       401  
   entertainment 386  
   Name: category, dtype: int64
```

```
sbn.countplot(df['category'])
```

```
↳ <matplotlib.axes._subplots.AxesSubplot at 0x7f449c30dc50>
```



```
# Use sklearn utility to convert label strings to numbered index  
from sklearn.preprocessing import LabelEncoder  
df["category"] = LabelEncoder().fit_transform(df["category"])  
df.head()
```



	category	text
0	4	tv future in the hands of viewers with home th...
1	0	worldcom boss left books alone former worldc...
2	3	tigers wary of farrell gamble leicester say ...
3	3	yeading face newcastle in fa cup premiership s...
4	1	ocean s twelve raids box office ocean s twelve...

```
#tokenize the words (text)
stemmer = PorterStemmer()
words = stopwords.words("english")
df['text'] = df['text'].apply(lambda x: " ".join([stemmer.stem(i) for i in re.sub("[^a-zA-Z]", " ", x).split() if i not in words]).lower())
vectorizer = TfidfVectorizer(min_df= 3, stop_words="english", sublinear_tf=True, norm='l2', ngram_range=(1, 2))
final_features = vectorizer.fit_transform(df['text']).toarray()
df.head(,)
```



	category	text
0	4	tv futur hand viewer home theatr system plasma...
1	0	worldcom boss left book alon former worldcom b...
2	3	tiger wari farrel gambl leicest say rush make ...
3	3	yead face newcastl fa cup premiership side new...
4	1	ocean twelv raid box offic ocean twelv crime c...

```
#split the data into training and test sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['category'], test_size=0.30, random_state=0)
# Inspect the dimenstions of our training and test data
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape,)
```



```
(1557,)
(668,)
(1557,)
(668,)
```

```
#convert to a vector with 1000 words
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(max_features = 1000)
vectorizer.fit(X_train)

X_train = vectorizer.transform(X_train)
X_test = vectorizer.transform(X_test)
X_train = X_train.todense()
```

```
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression()
classifier.fit(X_train, y_train)
score = classifier.score(X_test, y_test)
```

```
➤ /usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
  FutureWarning)
  /usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:469: FutureWarning: Default multi_class will be changed to '
    "this warning.", FutureWarning)
```

```
print("Accuracy of Logistic Regression:", score)
```

```
➤ Accuracy of Logistic Regression: 0.9745508982035929
```

```
# Converts the labels to a one-hot representation
import keras
num_classes = np.max(y_train) + 1
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
# Build the model
from tensorflow import keras
layers = keras.layers
models = keras.models
from keras.models import Sequential
from keras.layers import Dense
model = Sequential()

model.add(Dense(20, input_dim= 1000,activation='relu'))
model.add(Dense(20, activation='relu'))
model.add(Dense(num_classes , activation='softmax'))
```

```
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

```
model.summary()
```



Layer (type)	Output Shape	Param #
=====	=====	=====
dense_46 (Dense)	(None, 20)	20020
dense_47 (Dense)	(None, 20)	420
dense_48 (Dense)	(None, 5)	105
=====	=====	=====
Total params: 20,545		
Trainable params: 20,545		
Non-trainable params: 0		
=====		

```
history = model.fit(X_train, y_train,  
                    epochs=50,  
                    validation_split=0.10,  
                    batch_size=4)
```



Train on 1401 samples, validate on 156 samples

Epoch 1/50

1401/1401 [=====] - 2s 2ms/step - loss: 0.6100 - acc: 0.8144 - val_loss: 0.3113 - val_acc: 0.8974

Epoch 2/50

1401/1401 [=====] - 1s 901us/step - loss: 0.0733 - acc: 0.9843 - val_loss: 0.1114 - val_acc: 0.9808

Epoch 3/50

1401/1401 [=====] - 1s 899us/step - loss: 0.0187 - acc: 0.9979 - val_loss: 0.0808 - val_acc: 0.9808

Epoch 4/50

1401/1401 [=====] - 1s 884us/step - loss: 0.0068 - acc: 1.0000 - val_loss: 0.0778 - val_acc: 0.9744

Epoch 5/50

1401/1401 [=====] - 1s 929us/step - loss: 0.0039 - acc: 1.0000 - val_loss: 0.0814 - val_acc: 0.9744

Epoch 6/50

1401/1401 [=====] - 1s 905us/step - loss: 0.0024 - acc: 1.0000 - val_loss: 0.0764 - val_acc: 0.9744

Epoch 7/50

1401/1401 [=====] - 1s 910us/step - loss: 0.0016 - acc: 1.0000 - val_loss: 0.0794 - val_acc: 0.9744

Epoch 8/50

1401/1401 [=====] - 1s 868us/step - loss: 0.0011 - acc: 1.0000 - val_loss: 0.0784 - val_acc: 0.9744

Epoch 9/50

1401/1401 [=====] - 1s 873us/step - loss: 7.8040e-04 - acc: 1.0000 - val_loss: 0.0783 - val_acc: 0.9744

Epoch 10/50

1401/1401 [=====] - 1s 881us/step - loss: 5.6683e-04 - acc: 1.0000 - val_loss: 0.0801 - val_acc: 0.9744

Epoch 11/50

1401/1401 [=====] - 1s 880us/step - loss: 4.1964e-04 - acc: 1.0000 - val_loss: 0.0816 - val_acc: 0.9744

Epoch 12/50

1401/1401 [=====] - 1s 875us/step - loss: 3.1483e-04 - acc: 1.0000 - val_loss: 0.0842 - val_acc: 0.9744

Epoch 13/50

1401/1401 [=====] - 1s 878us/step - loss: 2.4039e-04 - acc: 1.0000 - val_loss: 0.0823 - val_acc: 0.9744

Epoch 14/50

1401/1401 [=====] - 1s 904us/step - loss: 1.8510e-04 - acc: 1.0000 - val_loss: 0.0858 - val_acc: 0.9744

Epoch 15/50

1401/1401 [=====] - 1s 930us/step - loss: 1.4411e-04 - acc: 1.0000 - val_loss: 0.0892 - val_acc: 0.9744

Epoch 16/50

1401/1401 [=====] - 1s 885us/step - loss: 1.1325e-04 - acc: 1.0000 - val_loss: 0.0900 - val_acc: 0.9744

Epoch 17/50

1401/1401 [=====] - 1s 901us/step - loss: 8.9144e-05 - acc: 1.0000 - val_loss: 0.0916 - val_acc: 0.9744

Epoch 18/50

1401/1401 [=====] - 1s 901us/step - loss: 7.0803e-05 - acc: 1.0000 - val_loss: 0.0950 - val_acc: 0.9744

Epoch 19/50

1401/1401 [=====] - 1s 923us/step - loss: 5.6671e-05 - acc: 1.0000 - val_loss: 0.0964 - val_acc: 0.9744

Epoch 20/50

1401/1401 [=====] - 1s 880us/step - loss: 4.5506e-05 - acc: 1.0000 - val_loss: 0.0963 - val_acc: 0.9744

Epoch 21/50

1401/1401 [=====] - 1s 879us/step - loss: 3.6594e-05 - acc: 1.0000 - val_loss: 0.0998 - val_acc: 0.9744

Epoch 22/50

1401/1401 [=====] - 1s 910us/step - loss: 2.9441e-05 - acc: 1.0000 - val_loss: 0.1000 - val_acc: 0.9744

Epoch 23/50

```
1401/1401 [=====] - 1s 874us/step - loss: 2.3763e-05 - acc: 1.0000 - val_loss: 0.1005 - val_acc: 0.9744
Epoch 24/50
1401/1401 [=====] - 1s 906us/step - loss: 1.9280e-05 - acc: 1.0000 - val_loss: 0.1030 - val_acc: 0.9744
Epoch 25/50
1401/1401 [=====] - 1s 881us/step - loss: 1.5571e-05 - acc: 1.0000 - val_loss: 0.1041 - val_acc: 0.9744
Epoch 26/50
1401/1401 [=====] - 1s 868us/step - loss: 1.2621e-05 - acc: 1.0000 - val_loss: 0.1061 - val_acc: 0.9744
Epoch 27/50
1401/1401 [=====] - 1s 872us/step - loss: 1.0231e-05 - acc: 1.0000 - val_loss: 0.1099 - val_acc: 0.9744
Epoch 28/50
1401/1401 [=====] - 1s 883us/step - loss: 8.3612e-06 - acc: 1.0000 - val_loss: 0.1122 - val_acc: 0.9744
Epoch 29/50
1401/1401 [=====] - 1s 873us/step - loss: 6.7757e-06 - acc: 1.0000 - val_loss: 0.1125 - val_acc: 0.9744
Epoch 30/50
1401/1401 [=====] - 1s 887us/step - loss: 5.5421e-06 - acc: 1.0000 - val_loss: 0.1159 - val_acc: 0.9744
Epoch 31/50
1401/1401 [=====] - 1s 890us/step - loss: 4.5160e-06 - acc: 1.0000 - val_loss: 0.1159 - val_acc: 0.9744
Epoch 32/50
1401/1401 [=====] - 1s 885us/step - loss: 3.6901e-06 - acc: 1.0000 - val_loss: 0.1183 - val_acc: 0.9744
Epoch 33/50
1401/1401 [=====] - 1s 895us/step - loss: 3.0240e-06 - acc: 1.0000 - val_loss: 0.1203 - val_acc: 0.9744
Epoch 34/50
1401/1401 [=====] - 1s 896us/step - loss: 2.4854e-06 - acc: 1.0000 - val_loss: 0.1217 - val_acc: 0.9744
Epoch 35/50
1401/1401 [=====] - 1s 882us/step - loss: 2.0442e-06 - acc: 1.0000 - val_loss: 0.1213 - val_acc: 0.9744
Epoch 36/50
1401/1401 [=====] - 1s 864us/step - loss: 1.6950e-06 - acc: 1.0000 - val_loss: 0.1263 - val_acc: 0.9744
Epoch 37/50
1401/1401 [=====] - 1s 869us/step - loss: 1.4025e-06 - acc: 1.0000 - val_loss: 0.1270 - val_acc: 0.9744
Epoch 38/50
1401/1401 [=====] - 1s 877us/step - loss: 1.1615e-06 - acc: 1.0000 - val_loss: 0.1271 - val_acc: 0.9744
Epoch 39/50
1401/1401 [=====] - 1s 883us/step - loss: 9.6589e-07 - acc: 1.0000 - val_loss: 0.1272 - val_acc: 0.9744
Epoch 40/50
1401/1401 [=====] - 1s 865us/step - loss: 8.0669e-07 - acc: 1.0000 - val_loss: 0.1300 - val_acc: 0.9744
Epoch 41/50
1401/1401 [=====] - 1s 845us/step - loss: 6.7676e-07 - acc: 1.0000 - val_loss: 0.1298 - val_acc: 0.9744
Epoch 42/50
1401/1401 [=====] - 1s 848us/step - loss: 5.7363e-07 - acc: 1.0000 - val_loss: 0.1328 - val_acc: 0.9744
Epoch 43/50
1401/1401 [=====] - 1s 846us/step - loss: 4.8343e-07 - acc: 1.0000 - val_loss: 0.1351 - val_acc: 0.9744
Epoch 44/50
1401/1401 [=====] - 1s 871us/step - loss: 4.1209e-07 - acc: 1.0000 - val_loss: 0.1389 - val_acc: 0.9744
Epoch 45/50
1401/1401 [=====] - 1s 854us/step - loss: 3.5461e-07 - acc: 1.0000 - val_loss: 0.1380 - val_acc: 0.9744
Epoch 46/50
1401/1401 [=====] - 1s 840us/step - loss: 3.0675e-07 - acc: 1.0000 - val_loss: 0.1410 - val_acc: 0.9744
```

```
Epoch 47/50
1401/1401 [=====] - 1s 846us/step - loss: 2.6828e-07 - acc: 1.0000 - val_loss: 0.1420 - val_acc: 0.9744
Epoch 48/50
1401/1401 [=====] - 1s 877us/step - loss: 2.3718e-07 - acc: 1.0000 - val_loss: 0.1445 - val_acc: 0.9744
Epoch 49/50
1401/1401 [=====] - 1s 853us/step - loss: 2.1174e-07 - acc: 1.0000 - val_loss: 0.1442 - val_acc: 0.9744
Epoch 50/50
1401/1401 [=====] - 1s 862us/step - loss: 1.9043e-07 - acc: 1.0000 - val_loss: 0.1465 - val_acc: 0.9744
```

```
loss, accuracy = model.evaluate(X_train, y_train, verbose=False)
print("Training Accuracy: {:.4f}".format(accuracy))
loss, accuracy = model.evaluate(X_test, y_test, verbose=False)
print("Testing Accuracy: {:.4f}".format(accuracy))
```

```
☞ Training Accuracy: 0.9974
   Testing Accuracy: 0.9775
```

```
import matplotlib.pyplot as plt
plt.style.use('ggplot')


def plot_history(history):
    acc = history.history['acc']
    val_acc = history.history['val_acc']
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    x = range(1, len(acc) + 1)

    plt.figure(figsize=(12, 5))
    plt.subplot(1, 2, 1)
    plt.plot(x, acc, 'b', label='Training acc')
    plt.plot(x, val_acc, 'r', label='Validation acc')
    plt.title('Training and validation accuracy')
    plt.legend()
    plt.subplot(1, 2, 2)
    plt.plot(x, loss, 'b', label='Training loss')
    plt.plot(x, val_loss, 'r', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend()
```

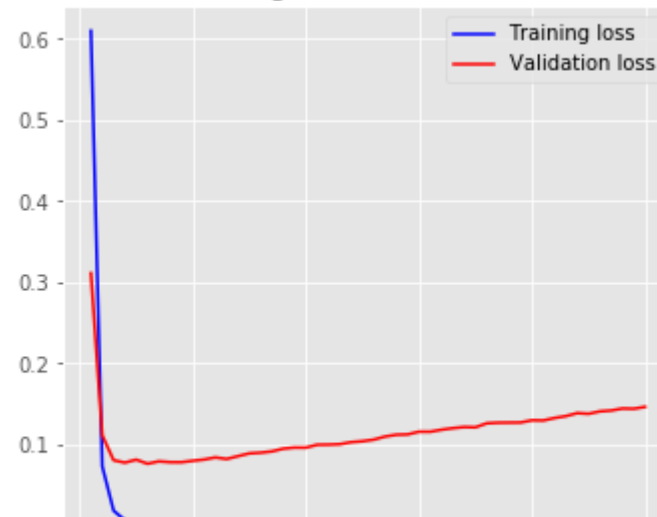
```
plot_history(history)
```

```
☞
```


Training and validation accuracy



Training and validation loss



```
y_softmax = model.predict(X_test)
```

```
y_test_1d = []
```

```
y_pred_1d = []
```

```
for i in range(len(y_test)):
    probs = y_test[i]
    index_arr = np.nonzero(probs)
    one_hot_index = index_arr[0].item(0)
    y_test_1d.append(one_hot_index)
```

```
for i in range(0, len(y_softmax)):
    probs = y_softmax[i]
    predicted_index = np.argmax(probs)
    y_pred_1d.append(predicted_index)
```

```
from sklearn import metrics
```

```
print(metrics.confusion_matrix(y_test_1d, y_pred_1d))
```

```
print(metrics.classification_report(y_test_1d, y_pred_1d))
```

```
from sklearn.metrics import accuracy_score
```

```
print("Accuracy of Deep Model is:", accuracy_score(y_test_1d, y_pred_1d))
```

